

Assign 6_final

December 6, 2017

```
In [50]: from tensorflow.examples.tutorials.mnist import input_data

import tensorflow as tf

from keras.models import Sequential
from keras.layers import Dense
from keras import initializers
import numpy as np
from numpy import newaxis
import matplotlib.pyplot as plt
import matplotlib.mlab as mlab
import keras as keras

import pandas as pd

In [61]: # 1A
from keras.optimizers import SGD
from keras.initializers import Zeros
from keras.datasets import mnist
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()

#####
print('Training data shape : ', train_images.shape, train_labels.shape)

print('Testing data shape : ', test_images.shape, test_labels.shape)

# Find the unique numbers from the train labels
classes = np.unique(train_labels)
nClasses = len(classes)
print('Total number of outputs : ', nClasses)
print('Output classes : ', classes)

# Change from matrix to array of dimension 28x28 to array of dimension 784
dimData = np.prod(train_images.shape[1:])
train_data = train_images.reshape(train_images.shape[0], dimData)
test_data = test_images.reshape(test_images.shape[0], dimData)
```

```

# Change to float datatype
train_data = train_data.astype('float32')
test_data = test_data.astype('float32')

# Scale the data to lie between 0 to 1
train_data /= 255
test_data /= 255

# Change the labels from integer to categorical data
train_labels_one_hot = to_categorical(train_labels)
test_labels_one_hot = to_categorical(test_labels)

model = Sequential()
model.add(Dense(10, activation='softmax', use_bias=True, kernel_initializer=Zeros(),
               bias_initializer=Zeros(), input_shape=(784,)))
model.summary()
model.compile(loss='categorical_crossentropy',
              optimizer=SGD(0.5),
              metrics=['accuracy'])

history = model.fit(train_data, train_labels_one_hot, batch_size=100, epochs=17, verbose=0,
                    validation_data=(test_data, test_labels_one_hot),)

plt.gcf().clear()
plt.figure(1)

# summarize history for accuracy

plt.subplot(211)
plt.plot(history.history['acc'])
plt.plot(history.history['val_acc'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')

# summarize history for loss

plt.subplot(212)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()


```

Training data shape : (60000, 28, 28) (60000,)

Testing data shape : (10000, 28, 28) (10000,)
 Total number of outputs : 10
 Output classes : [0 1 2 3 4 5 6 7 8 9]

Layer (type)	Output Shape	Param #
dense_105 (Dense)	(None, 10)	7850

Total params: 7,850
 Trainable params: 7,850
 Non-trainable params: 0



```
In [52]: # 1B
# now in Keras (since it's easier to use
# https://www.learnopencv.com/image-classification-using-feedforward-neural-network-i

from keras.datasets import mnist
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()

from keras.utils import to_categorical

# Find the unique numbers from the train labels
```

```

classes = np.unique(train_labels)
nClasses = len(classes)

# Change from matrix to array of dimension 28x28 to array of dimension 784
dimData = np.prod(train_images.shape[1:])
train_data = train_images.reshape(train_images.shape[0], dimData)
test_data = test_images.reshape(test_images.shape[0], dimData)

# Change to float datatype
train_data = train_data.astype('float32')
test_data = test_data.astype('float32')

# Scale the data to lie between 0 to 1
train_data /= 255
test_data /= 255

# Change the labels from integer to categorical data
train_labels_one_hot = to_categorical(train_labels)
test_labels_one_hot = to_categorical(test_labels)

#-----
#weights
weight1 = TruncatedNormal(mean=0.0, stddev=0.01)
weight1((784,1500))
weight2 = TruncatedNormal(mean=0.0, stddev=0.01)
weight2((1500,1500))
weight3 = TruncatedNormal(mean=0.0, stddev=0.01)
weight3((1500, 1500))
weight4 = TruncatedNormal(mean=0.0, stddev=0.01)
weight4((1500, 10))

#Network Creation
model = Sequential()
model.add(Dense(1500, activation='relu',use_bias=True, input_shape=(dimData,), kernel_initializer=weight1))
model.add(Dense(1500, activation='relu',use_bias=True, kernel_initializer=weight2, bias_initializer=weight4))
model.add(Dense(1500, activation='relu',use_bias=True, kernel_initializer = weight3, bias_initializer=weight4))
model.add(Dense(nClasses, activation='softmax',use_bias=True, kernel_initializer=weight4))
model.summary()

#implementing adam optimizer
adam = keras.optimizers.Adam(lr=0.001, beta_1=0.9, beta_2=0.999, epsilon=1*10**(-8), clipvalue=1)

model.compile(optimizer=adam, loss='categorical_crossentropy', metrics=['accuracy'])

history = model.fit(train_data, train_labels_one_hot, batch_size=100, epochs=34, verbose=1,
                    validation_data=(test_data, test_labels_one_hot))

```

```

[test_loss, test_acc] = model.evaluate(test_data, test_labels_one_hot)

print(history.history.keys())
plt.gcf().clear()
plt.figure(1)

# summarize history for accuracy
plt.subplot(211)
plt.plot(history.history['acc'])
plt.plot(history.history['val_acc'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')

# summarize history for loss

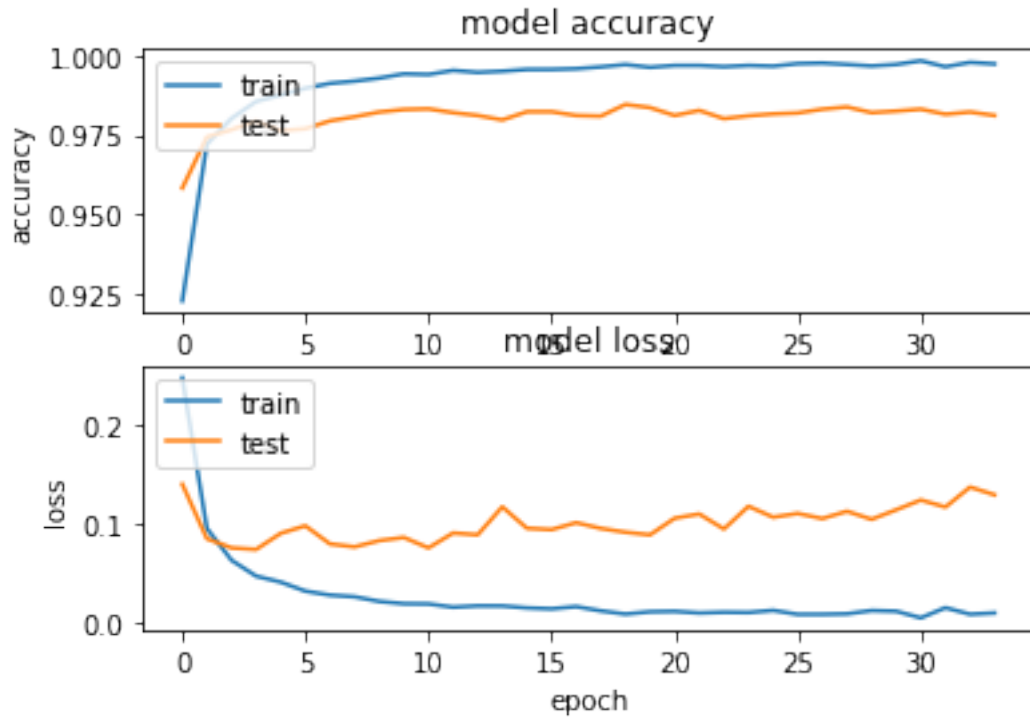
plt.subplot(212)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()

```

```

-----
Layer (type)                Output Shape                Param #
=====
dense_88 (Dense)             (None, 1500)                1177500
-----
dense_89 (Dense)             (None, 1500)                2251500
-----
dense_90 (Dense)             (None, 1500)                2251500
-----
dense_91 (Dense)             (None, 10)                  15010
=====
Total params: 5,695,510
Trainable params: 5,695,510
Non-trainable params: 0
-----
10000/10000 [=====] - 3s 261us/step
dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])

```



```
In [55]: # 1C
         # Dropout is the only real addition
         # Our model will contain the dropout() statement

from keras.layers import Dropout
from keras.datasets import mnist

(train_images, train_labels), (test_images, test_labels) = mnist.load_data()

from keras.utils import to_categorical

print('Training data shape : ', train_images.shape, train_labels.shape)

print('Testing data shape : ', test_images.shape, test_labels.shape)

# Find the unique numbers from the train labels
classes = np.unique(train_labels)
nClasses = len(classes)
print('Total number of outputs : ', nClasses)
print('Output classes : ', classes)

# Change from matrix to array of dimension 28x28 to array of dimension 784
dimData = np.prod(train_images.shape[1:])
```

```

train_data = train_images.reshape(train_images.shape[0], dimData)
test_data = test_images.reshape(test_images.shape[0], dimData)

# Change to float datatype
train_data = train_data.astype('float32')
test_data = test_data.astype('float32')

# Scale the data to lie between 0 to 1
train_data /= 255
test_data /= 255

# Change the labels from integer to categorical data
train_labels_one_hot = to_categorical(train_labels)
test_labels_one_hot = to_categorical(test_labels)

# Display the change for category label using one-hot encoding
print('Original label 0 : ', train_labels[0])
print('After conversion to categorical ( one-hot ) : ', train_labels_one_hot[0])

#-----
weight1 = TruncatedNormal(mean=0.0, stddev=0.01)
weight1((784,1500))
weight2 = TruncatedNormal(mean=0.0, stddev=0.01)
weight2((1500,1500))
weight3 = TruncatedNormal(mean=0.0, stddev=0.01)
weight3((1500, 1500))
weight4 = TruncatedNormal(mean=0.0, stddev=0.01)
weight4((1500, 10))

model = Sequential()
model.add(Dense(1500, activation='relu',use_bias=True, input_shape=(dimData,), kernel_initializer=weight1))
model.add(Dropout(0.5))
model.add(Dense(1500, activation='relu',use_bias=True, kernel_initializer=weight2, bias_initializer=weight4))
model.add(Dropout(0.5))
model.add(Dense(1500, activation='relu',use_bias=True, kernel_initializer = weight3, bias_initializer=weight4))
model.add(Dropout(0.5))
model.add(Dense(nClasses, activation='softmax',use_bias=True, kernel_initializer=weight4))

model.summary()

#implementing adam optimizer
adam = keras.optimizers.Adam(lr=0.001, beta_1=0.9, beta_2=0.999, epsilon=1*10**(-8), clipvalue=1)

model.compile(optimizer=adam, loss='categorical_crossentropy', metrics=['accuracy'])

history = model.fit(train_data, train_labels_one_hot, batch_size=100, epochs=34, verbose=1)

```

```

validation_data=(test_data, test_labels_one_hot),)

[test_loss, test_acc] = model.evaluate(test_data, test_labels_one_hot)
#print("Evaluation result on Test Data : Loss = {}, accuracy = {}".format(test_loss,
plt.gcf().clear()
plt.figure(1)

# summarize history for accuracy

plt.subplot(211)
plt.plot(history.history['acc'])
plt.plot(history.history['val_acc'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')

# summarize history for loss

plt.subplot(212)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()

```

Training data shape : (60000, 28, 28) (60000,)

Testing data shape : (10000, 28, 28) (10000,)

Total number of outputs : 10

Output classes : [0 1 2 3 4 5 6 7 8 9]

Original label 0 : 5

After conversion to categorical (one-hot) : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]

Layer (type)	Output Shape	Param #
dense_96 (Dense)	(None, 1500)	1177500
dropout_46 (Dropout)	(None, 1500)	0
dense_97 (Dense)	(None, 1500)	2251500
dropout_47 (Dropout)	(None, 1500)	0
dense_98 (Dense)	(None, 1500)	2251500

dropout_48 (Dropout)	(None, 1500)	0
----------------------	--------------	---

dense_99 (Dense)	(None, 10)	15010
------------------	------------	-------

Total params: 5,695,510
 Trainable params: 5,695,510
 Non-trainable params: 0

Train on 60000 samples, validate on 10000 samples

Epoch 1/34

60000/60000 [=====] - 62s 1ms/step - loss: 0.3061 - acc: 0.9038 - val_

Epoch 2/34

60000/60000 [=====] - 59s 990us/step - loss: 0.1556 - acc: 0.9543 - va

Epoch 3/34

60000/60000 [=====] - 60s 1ms/step - loss: 0.1280 - acc: 0.9627 - val_

Epoch 4/34

60000/60000 [=====] - 59s 985us/step - loss: 0.1172 - acc: 0.9658 - va

Epoch 5/34

60000/60000 [=====] - 59s 985us/step - loss: 0.1041 - acc: 0.9702 - va

Epoch 6/34

60000/60000 [=====] - 59s 986us/step - loss: 0.0997 - acc: 0.9714 - va

Epoch 7/34

60000/60000 [=====] - 59s 984us/step - loss: 0.0935 - acc: 0.9745 - va

Epoch 8/34

60000/60000 [=====] - 59s 989us/step - loss: 0.0851 - acc: 0.9759 - va

Epoch 9/34

60000/60000 [=====] - 59s 982us/step - loss: 0.0863 - acc: 0.9761 - va

Epoch 10/34

60000/60000 [=====] - 59s 981us/step - loss: 0.0820 - acc: 0.9774 - va

Epoch 11/34

60000/60000 [=====] - 59s 980us/step - loss: 0.0805 - acc: 0.9783 - va

Epoch 12/34

60000/60000 [=====] - 59s 979us/step - loss: 0.0770 - acc: 0.9795 - va

Epoch 13/34

60000/60000 [=====] - 59s 980us/step - loss: 0.0769 - acc: 0.9797 - va

Epoch 14/34

60000/60000 [=====] - 59s 979us/step - loss: 0.0707 - acc: 0.9802 - va

Epoch 15/34

60000/60000 [=====] - 59s 978us/step - loss: 0.0695 - acc: 0.9813 - va

Epoch 16/34

60000/60000 [=====] - 59s 978us/step - loss: 0.0698 - acc: 0.9820 - va

Epoch 17/34

60000/60000 [=====] - 59s 977us/step - loss: 0.0704 - acc: 0.9814 - va

Epoch 18/34

60000/60000 [=====] - 59s 983us/step - loss: 0.0694 - acc: 0.9819 - va

Epoch 19/34


60000/60000 [=====] - 60s 994us/step - loss: 0.0691 - acc: 0.9820 - va

Epoch 20/34

```

60000/60000 [=====] - 59s 988us/step - loss: 0.0714 - acc: 0.9817 - va
Epoch 21/34
60000/60000 [=====] - 60s 1000us/step - loss: 0.0646 - acc: 0.9833 - v
Epoch 22/34
60000/60000 [=====] - 60s 998us/step - loss: 0.0656 - acc: 0.9841 - va
Epoch 23/34
60000/60000 [=====] - 60s 992us/step - loss: 0.0648 - acc: 0.9846 - va
Epoch 24/34
60000/60000 [=====] - 60s 997us/step - loss: 0.0708 - acc: 0.9831 - va
Epoch 25/34
60000/60000 [=====] - 60s 998us/step - loss: 0.0642 - acc: 0.9844 - va
Epoch 26/34
60000/60000 [=====] - 60s 999us/step - loss: 0.0664 - acc: 0.9839 - va
Epoch 27/34
60000/60000 [=====] - 60s 994us/step - loss: 0.0630 - acc: 0.9844 - va
Epoch 28/34
60000/60000 [=====] - 60s 996us/step - loss: 0.0685 - acc: 0.9844 - va
Epoch 29/34
60000/60000 [=====] - 60s 997us/step - loss: 0.0649 - acc: 0.9852 - va
Epoch 30/34
60000/60000 [=====] - 60s 1ms/step - loss: 0.0687 - acc: 0.9844 - val
Epoch 31/34
60000/60000 [=====] - 60s 1ms/step - loss: 0.0673 - acc: 0.9856 - val
Epoch 32/34
60000/60000 [=====] - 60s 1ms/step - loss: 0.0649 - acc: 0.9863 - val
Epoch 33/34
60000/60000 [=====] - 60s 999us/step - loss: 0.0657 - acc: 0.9854 - va
Epoch 34/34
60000/60000 [=====] - 60s 999us/step - loss: 0.0737 - acc: 0.9842 - va
10000/10000 [=====] - 3s 262us/step

```



In [58]: # 1D

```
from numpy.random import seed
seed(1)
from tensorflow import set_random_seed
set_random_seed(1)
import keras
import matplotlib.pyplot as plt
import numpy as np
from keras.datasets import mnist
from keras.models import Sequential
from keras.initializers import TruncatedNormal, Constant, Zeros
from keras.layers import Dense, Dropout, Conv2D, MaxPooling2D, Flatten
from keras.optimizers import SGD, Adam
# setting up of batch, and the number of classes and epochs

batch_size = 100
num_classes = 10
epochs = 34

# input image dimensions
img_x, img_y = 28, 28

# load the MNIST data set, which already splits into train and test sets for us
```

```

(x_train, y_train), (x_test, y_test) = mnist.load_data()

# reshape the data into a 4D tensor - (sample_number, x_img_size, y_img_size, num_channels)
# because the MNIST is greyscale, we only have a single channel - RGB colour images w
x_train = x_train.reshape(x_train.shape[0], img_x, img_y, 1)
x_test = x_test.reshape(x_test.shape[0], img_x, img_y, 1)
input_shape = (img_x, img_y, 1)

# convert the data to the right type
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

# convert class vectors to binary class matrices - this is for use in the
# categorical_crossentropy loss below
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

adam = keras.optimizers.Adam(lr=0.001, beta_1=0.9, beta_2=0.999, epsilon=1*10**(-8), clipvalue=1)

model = Sequential()
model.add(Conv2D(32,(5,5),strides=(1, 1),padding='same', activation='relu',use_bias=True))
model.add(MaxPooling2D(pool_size=(2, 2), strides=(2,2)))
model.add(Conv2D(64,(5,5),strides=(1, 1),padding='same', activation='relu',use_bias=True))
model.add(MaxPooling2D(pool_size=(2, 2), strides=(2,2)))
model.add(Flatten())
model.add(Dense(num_classes, activation='softmax',use_bias=True, bias_initializer=Constant(0)))

model.summary()

model.compile(loss=keras.losses.categorical_crossentropy,optimizer=adam,metrics=['accuracy'])

history = model.fit(x_train, y_train,batch_size=batch_size,epochs=epochs,verbose=0,validation_data=(x_test, y_test))

#[test_loss, test_acc] = model.evaluate(test_data, test_labels_one_hot)
#print("Evaluation result on Test Data : Loss = {}, accuracy = {}".format(test_loss, test_acc))
plt.gcf().clear()
plt.figure(1)

# summarize history for accuracy

plt.subplot(211)
plt.plot(history.history['acc'])

```

```
plt.plot(history.history['val_acc'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
```


summarize history for loss

```
plt.subplot(212)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```

x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples

Layer (type)	Output Shape	Param #
conv2d_27 (Conv2D)	(None, 28, 28, 32)	832
max_pooling2d_27 (MaxPooling)	(None, 14, 14, 32)	0
conv2d_28 (Conv2D)	(None, 14, 14, 64)	51264
max_pooling2d_28 (MaxPooling)	(None, 7, 7, 64)	0
flatten_14 (Flatten)	(None, 3136)	0
dense_102 (Dense)	(None, 10)	31370

Total params: 83,466
Trainable params: 83,466
Non-trainable params: 0



```
In [ ]: #1e
```

Plots of the subtasks were shown below of every task.

```
In [62]: from keras.layers import LSTM
```

```
#A
train_size = 8000
# np.random.seed(7)
dataset = np.random.randint(low=0, high=9+1, size=(10000,30))

def get_labels(dataset):
    sum = np.sum(dataset,1)
    sum[sum<100]=0
    sum[sum>=100]=1
    return sum

train, test = dataset[0:train_size,:], dataset[train_size:len(dataset),:]
train_labels = np.reshape(get_labels(train),(-1,1))
test_labels = np.reshape(get_labels(test),(-1,1))
train = train[:, :, newaxis]
test = test[:, :, newaxis]
# print(train.shape)
```

```

# print(train.ndim)
# print(len(train), len(test))
#B
timesteps=30
# create and fit the LSTM network
model = Sequential()
model.add(LSTM(200, input_shape=(train.shape[1:])))
model.add(Dense(1, activation='sigmoid'))
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
print(model.summary())

#D change epochs to 60
model.fit(train, train_labels, validation_data=(test, test_labels), epochs=60, batch_size=32)

#E
# Final evaluation of the model
scores = model.evaluate(test, test_labels, verbose=0)
print("Accuracy: %.2f%%" % (scores[1]*100))

```

```

-----
Layer (type)                 Output Shape              Param #
-----
lstm_1 (LSTM)                (None, 200)              161600
-----
dense_106 (Dense)            (None, 1)                201
-----
Total params: 161,801
Trainable params: 161,801
Non-trainable params: 0
-----
None
Train on 8000 samples, validate on 2000 samples
Epoch 1/60
8000/8000 [=====] - 10s 1ms/step - loss: 0.0655 - acc: 0.9856 - val_loss: 0.0478
Epoch 2/60
8000/8000 [=====] - 8s 991us/step - loss: 0.0478 - acc: 0.9866 - val_loss: 0.0379
Epoch 3/60
8000/8000 [=====] - 8s 1ms/step - loss: 0.0379 - acc: 0.9874 - val_loss: 0.0343
Epoch 4/60
8000/8000 [=====] - 8s 1ms/step - loss: 0.0343 - acc: 0.9888 - val_loss: 0.0261
Epoch 5/60
8000/8000 [=====] - 8s 1ms/step - loss: 0.0261 - acc: 0.9914 - val_loss: 0.0249
Epoch 6/60
8000/8000 [=====] - 8s 1ms/step - loss: 0.0249 - acc: 0.9908 - val_loss: 0.0190
Epoch 7/60
8000/8000 [=====] - 8s 1ms/step - loss: 0.0190 - acc: 0.9924 - val_loss: 0.0197
Epoch 8/60
8000/8000 [=====] - 8s 1ms/step - loss: 0.0197 - acc: 0.9934 - val_loss: 0.0197

```

Epoch 9/60
8000/8000 [=====] - 8s 1ms/step - loss: 0.0181 - acc: 0.9931 - val_loss: 0.0181
Epoch 10/60
8000/8000 [=====] - 8s 1ms/step - loss: 0.0204 - acc: 0.9929 - val_loss: 0.0204
Epoch 11/60
8000/8000 [=====] - 8s 1ms/step - loss: 0.0163 - acc: 0.9936 - val_loss: 0.0163
Epoch 12/60
8000/8000 [=====] - 8s 1ms/step - loss: 0.0122 - acc: 0.9959 - val_loss: 0.0122
Epoch 13/60
8000/8000 [=====] - 8s 1ms/step - loss: 0.0142 - acc: 0.9946 - val_loss: 0.0142
Epoch 14/60
8000/8000 [=====] - 8s 1ms/step - loss: 0.0127 - acc: 0.9948 - val_loss: 0.0127
Epoch 15/60
8000/8000 [=====] - 8s 994us/step - loss: 0.0156 - acc: 0.9928 - val_loss: 0.0156
Epoch 16/60
8000/8000 [=====] - 8s 995us/step - loss: 0.0130 - acc: 0.9954 - val_loss: 0.0130
Epoch 17/60
8000/8000 [=====] - 8s 991us/step - loss: 0.0111 - acc: 0.9960 - val_loss: 0.0111
Epoch 18/60
8000/8000 [=====] - 8s 982us/step - loss: 0.0120 - acc: 0.9958 - val_loss: 0.0120
Epoch 19/60
8000/8000 [=====] - 8s 989us/step - loss: 0.0115 - acc: 0.9954 - val_loss: 0.0115
Epoch 20/60
8000/8000 [=====] - 8s 991us/step - loss: 0.0118 - acc: 0.9953 - val_loss: 0.0118
Epoch 21/60
8000/8000 [=====] - 8s 994us/step - loss: 0.0113 - acc: 0.9950 - val_loss: 0.0113
Epoch 22/60
8000/8000 [=====] - 8s 998us/step - loss: 0.0101 - acc: 0.9960 - val_loss: 0.0101
Epoch 23/60
8000/8000 [=====] - 8s 1ms/step - loss: 0.0110 - acc: 0.9956 - val_loss: 0.0110
Epoch 24/60
8000/8000 [=====] - 8s 989us/step - loss: 0.0122 - acc: 0.9953 - val_loss: 0.0122
Epoch 25/60
8000/8000 [=====] - 8s 996us/step - loss: 0.0084 - acc: 0.9973 - val_loss: 0.0084
Epoch 26/60
8000/8000 [=====] - 8s 991us/step - loss: 0.0123 - acc: 0.9951 - val_loss: 0.0123
Epoch 27/60
8000/8000 [=====] - 8s 991us/step - loss: 0.0092 - acc: 0.9965 - val_loss: 0.0092
Epoch 28/60
8000/8000 [=====] - 8s 994us/step - loss: 0.0104 - acc: 0.9955 - val_loss: 0.0104
Epoch 29/60
8000/8000 [=====] - 8s 991us/step - loss: 0.0111 - acc: 0.9964 - val_loss: 0.0111
Epoch 30/60
8000/8000 [=====] - 8s 987us/step - loss: 0.0096 - acc: 0.9956 - val_loss: 0.0096
Epoch 31/60
8000/8000 [=====] - 8s 1ms/step - loss: 0.0071 - acc: 0.9975 - val_loss: 0.0071
Epoch 32/60
8000/8000 [=====] - 8s 1ms/step - loss: 0.0083 - acc: 0.9968 - val_loss: 0.0083

Epoch 33/60
8000/8000 [=====] - 8s 1ms/step - loss: 0.0101 - acc: 0.9964 - val_loss: 0.0101
Epoch 34/60
8000/8000 [=====] - 8s 1ms/step - loss: 0.0070 - acc: 0.9976 - val_loss: 0.0070
Epoch 35/60
8000/8000 [=====] - 8s 1ms/step - loss: 0.0065 - acc: 0.9975 - val_loss: 0.0065
Epoch 36/60
8000/8000 [=====] - 8s 1ms/step - loss: 0.0105 - acc: 0.9959 - val_loss: 0.0105
Epoch 37/60
8000/8000 [=====] - 8s 992us/step - loss: 0.0066 - acc: 0.9971 - val_loss: 0.0066
Epoch 38/60
8000/8000 [=====] - 8s 1ms/step - loss: 0.0083 - acc: 0.9974 - val_loss: 0.0083
Epoch 39/60
8000/8000 [=====] - 8s 992us/step - loss: 0.0080 - acc: 0.9965 - val_loss: 0.0080
Epoch 40/60
8000/8000 [=====] - 8s 993us/step - loss: 0.0087 - acc: 0.9966 - val_loss: 0.0087
Epoch 41/60
8000/8000 [=====] - 8s 993us/step - loss: 0.0061 - acc: 0.9980 - val_loss: 0.0061
Epoch 42/60
8000/8000 [=====] - 8s 1ms/step - loss: 0.0062 - acc: 0.9980 - val_loss: 0.0062
Epoch 43/60
8000/8000 [=====] - 8s 1ms/step - loss: 0.0069 - acc: 0.9970 - val_loss: 0.0069
Epoch 44/60
8000/8000 [=====] - 8s 1ms/step - loss: 0.0085 - acc: 0.9973 - val_loss: 0.0085
Epoch 45/60
8000/8000 [=====] - 8s 1ms/step - loss: 0.0060 - acc: 0.9976 - val_loss: 0.0060
Epoch 46/60
8000/8000 [=====] - 8s 1ms/step - loss: 0.0090 - acc: 0.9965 - val_loss: 0.0090
Epoch 47/60
8000/8000 [=====] - 8s 998us/step - loss: 0.0056 - acc: 0.9981 - val_loss: 0.0056
Epoch 48/60
8000/8000 [=====] - 8s 991us/step - loss: 0.0064 - acc: 0.9979 - val_loss: 0.0064
Epoch 49/60
8000/8000 [=====] - 8s 993us/step - loss: 0.0065 - acc: 0.9971 - val_loss: 0.0065
Epoch 50/60
8000/8000 [=====] - 8s 995us/step - loss: 0.0047 - acc: 0.9981 - val_loss: 0.0047
Epoch 51/60
8000/8000 [=====] - 8s 995us/step - loss: 0.0061 - acc: 0.9975 - val_loss: 0.0061
Epoch 52/60
8000/8000 [=====] - 8s 996us/step - loss: 0.0100 - acc: 0.9968 - val_loss: 0.0100
Epoch 53/60
8000/8000 [=====] - 8s 1ms/step - loss: 0.0051 - acc: 0.9981 - val_loss: 0.0051
Epoch 54/60
8000/8000 [=====] - 8s 1ms/step - loss: 0.0050 - acc: 0.9981 - val_loss: 0.0050
Epoch 55/60
8000/8000 [=====] - 8s 1ms/step - loss: 0.0059 - acc: 0.9975 - val_loss: 0.0059
Epoch 56/60
8000/8000 [=====] - 9s 1ms/step - loss: 0.0061 - acc: 0.9975 - val_loss: 0.0061

```
Epoch 57/60
8000/8000 [=====] - 8s 1ms/step - loss: 0.0040 - acc: 0.9988 - val_loss: 0.0040
Epoch 58/60
8000/8000 [=====] - 8s 991us/step - loss: 0.0035 - acc: 0.9991 - val_loss: 0.0035
Epoch 59/60
8000/8000 [=====] - 8s 994us/step - loss: 0.0047 - acc: 0.9983 - val_loss: 0.0047
Epoch 60/60
8000/8000 [=====] - 8s 991us/step - loss: 0.0023 - acc: 0.9993 - val_loss: 0.0023
Accuracy: 99.65%
```